

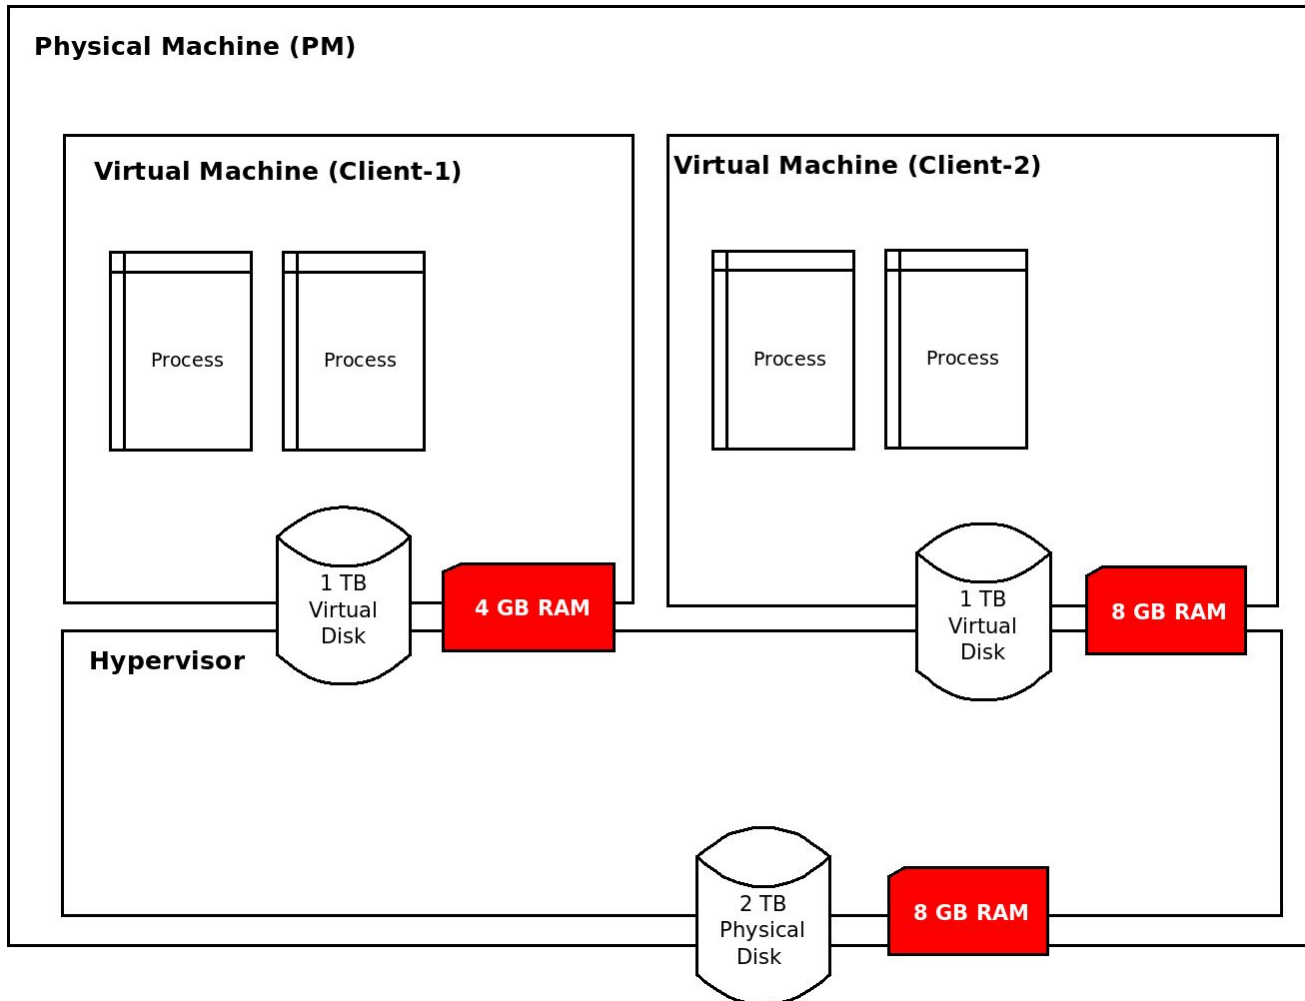
Multilevel Differentiated Hypervisor Caching for Derivative Clouds

Sprint Thesis Talk, RISC '17

Prashanth, M.Tech. II

Advisor:
Prof. Purushottam Kulkarni

Cloud Provider Architecture

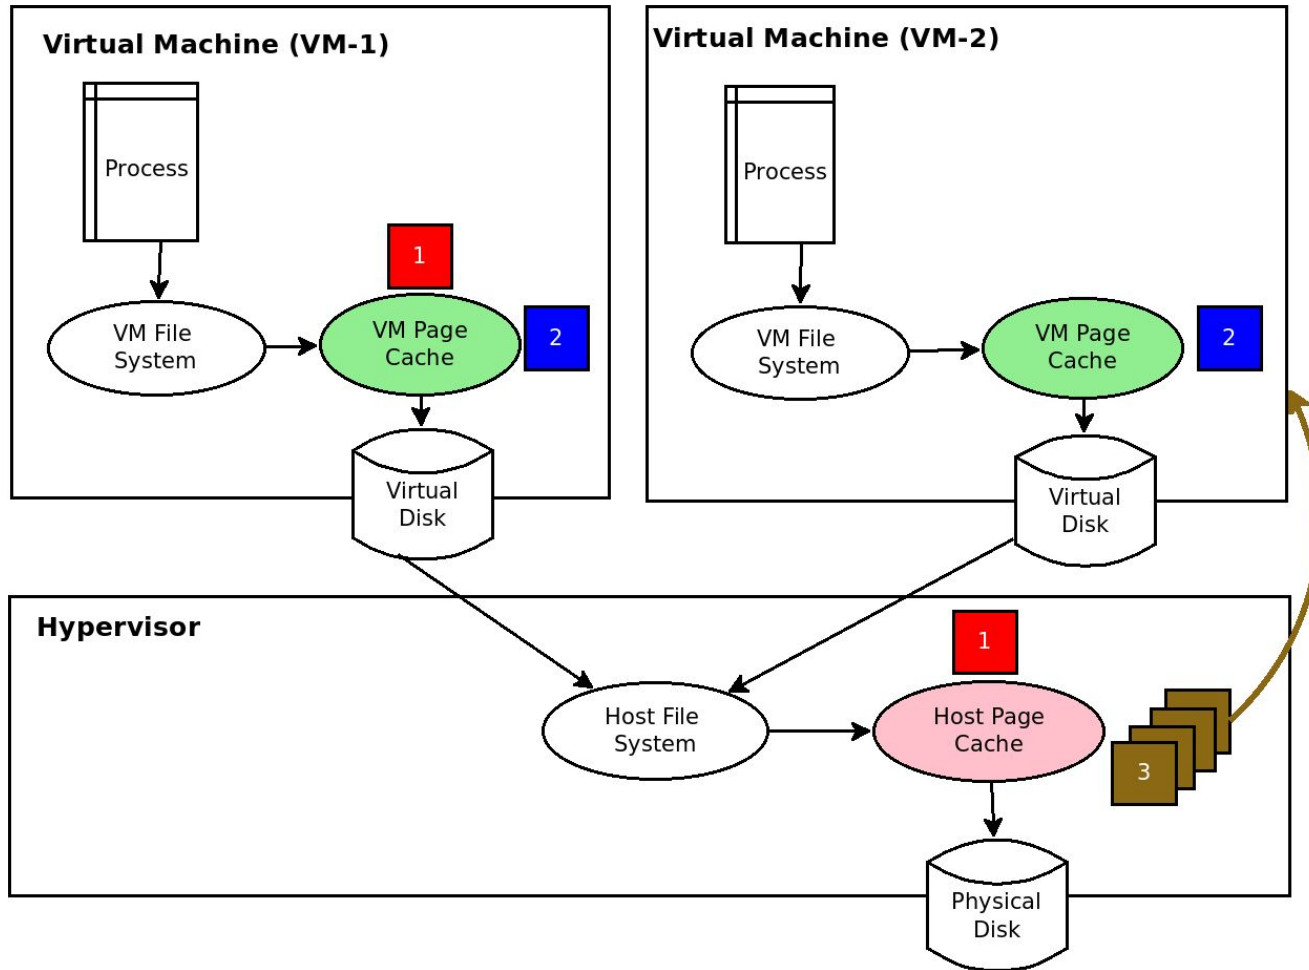


Key Points

1. Provision clients on VMs
2. Map SLA requirements to VM resources
3. Overprovision resources for cost benefits

Fig-1: Cloud providers provisioning clients using VMs

Caching in a cloud setup

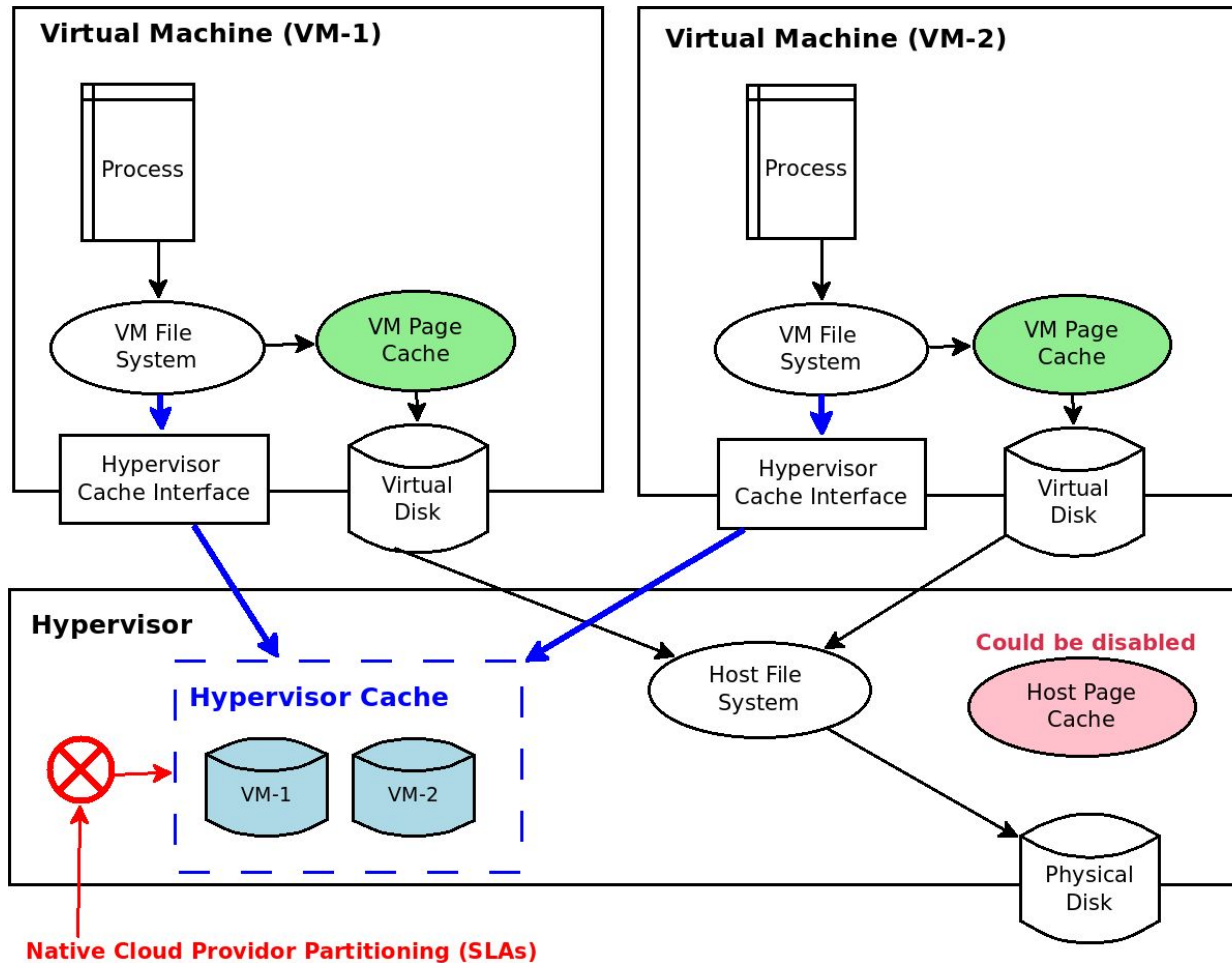


Drawbacks

1. Multiple copies of same page at host and VM
2. Multiple copies between VMs
3. Flooding of host page cache by a particular VM

Fig-2: File I/O in traditional hypervisor based cloud setup

Hypervisor managed caching



Key Points

1. Partitioning per VM based on high level SLAs
2. Exclusive cache
3. Dynamic readjustments
4. Existing works address this issue -
[SDC SoCC '15]
[Centaur ICAC '15]

Fig-3: Hypervisor Caching

 Marks the controller

What are traditional caches backed by ?

- ❑ Caches could be backed by
 1. Memory (RAM)
 2. SSD
 3. NVMs etc.
- ❑ We could even combined them to form multi-level or hybrid cache designs - [ExTmem HPCC '14]
- ❑ Existing literatures on hypervisor caches are of single level cache

Derivative Cloud Environment

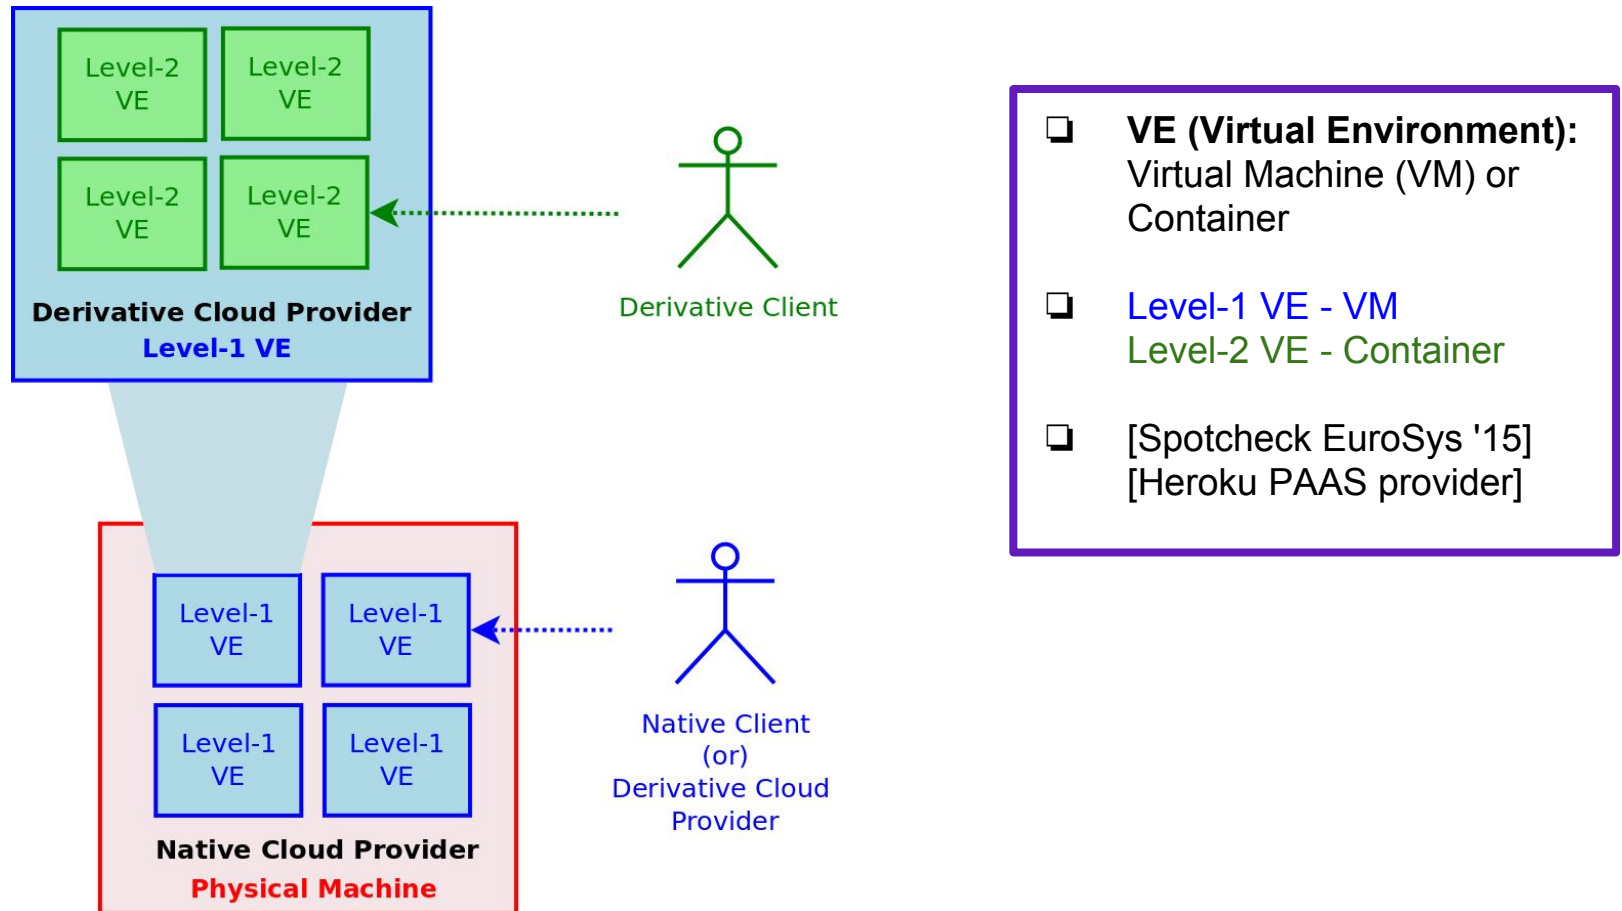


Fig-4: Comparison between native and derivative cloud environment

Problem Statement

To develop a caching framework that supports,

- ❑ Hypervisor caching
- ❑ Multiple levels of configurable cache
- ❑ A derivative cloud setup for enforce native and derivative provider SLA policies

Requirements of desired system

- ❑ Multiple levels of hypervisor managed caches
- ❑ Per VM and per container configurable caches at each level
- ❑ Resource conserving nature
- ❑ Spillover mechanism - Exceeds in L1 spilled over to the L2 cache
- ❑ Exclusive caching at all levels

Hypervisor caching in derivative clouds

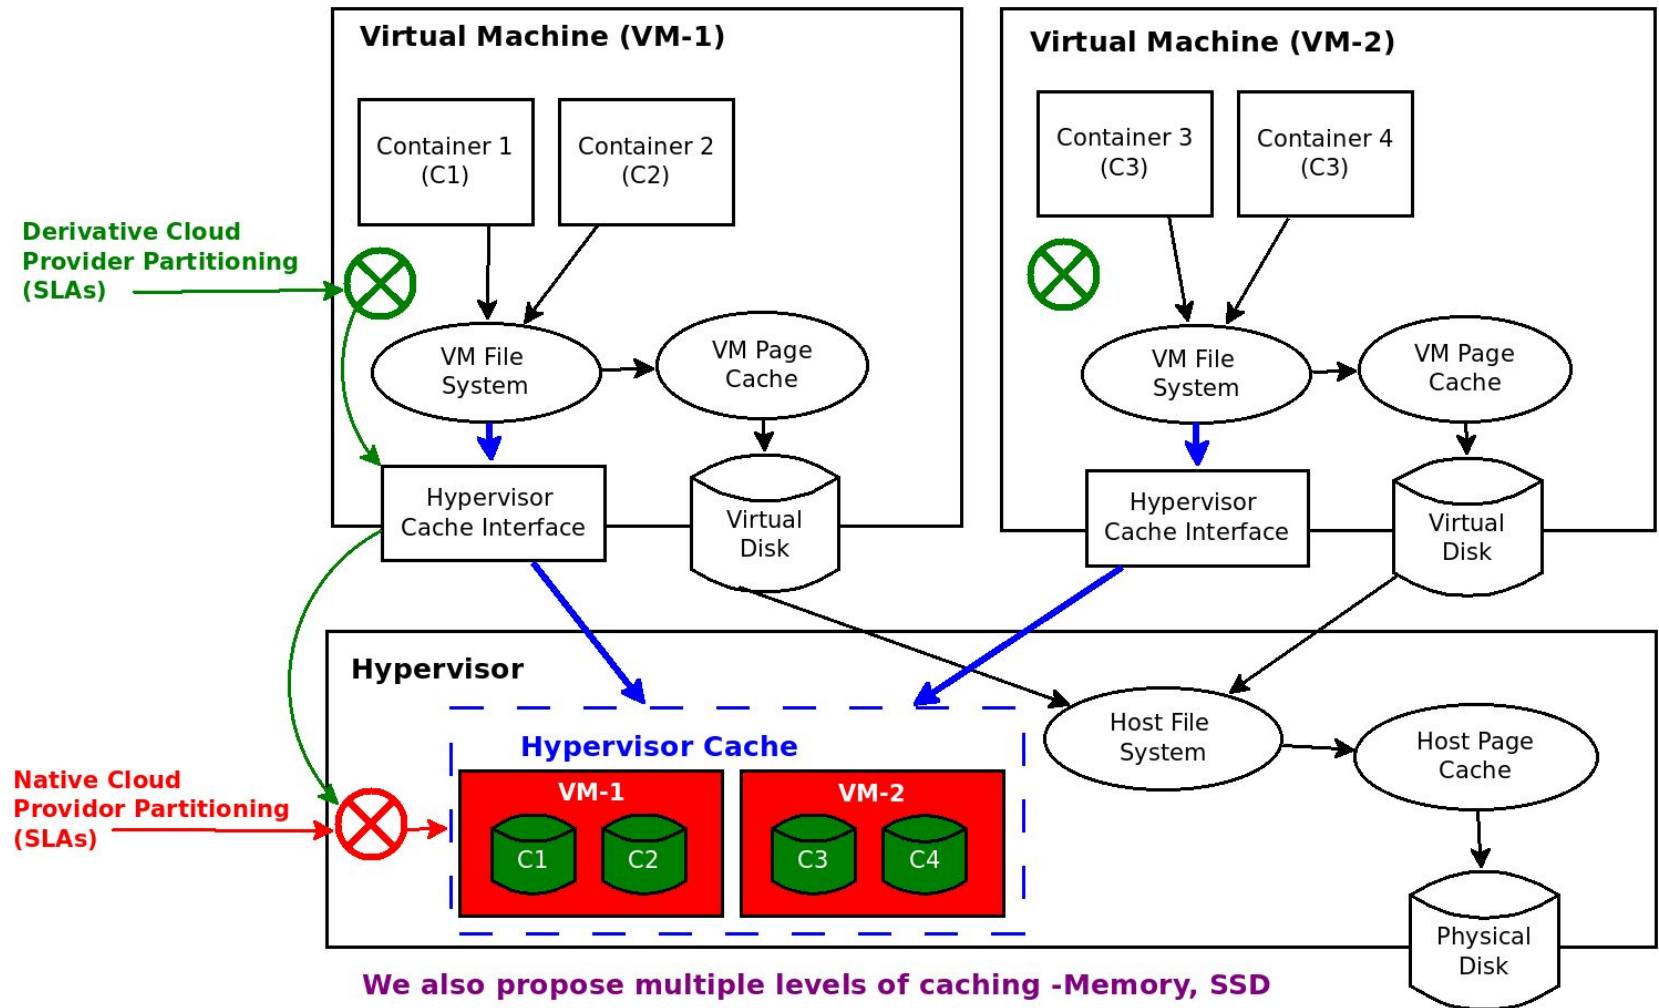


Fig-5: Proposed Architecture for cache partitioning in derivative clouds

Implementation Specifics

- ❑ T-MEM (Transcendent Memory) cache - A second cache caching framework for optimization of RAM
- ❑ Extended this to support hypervisor backed caches using memory and SSD
- ❑ KVM Hypervisor
- ❑ LXC containers (could be easily extended to other container managers)
- ❑ Control Knobs - Relative weights, Cache size

Current Implementation

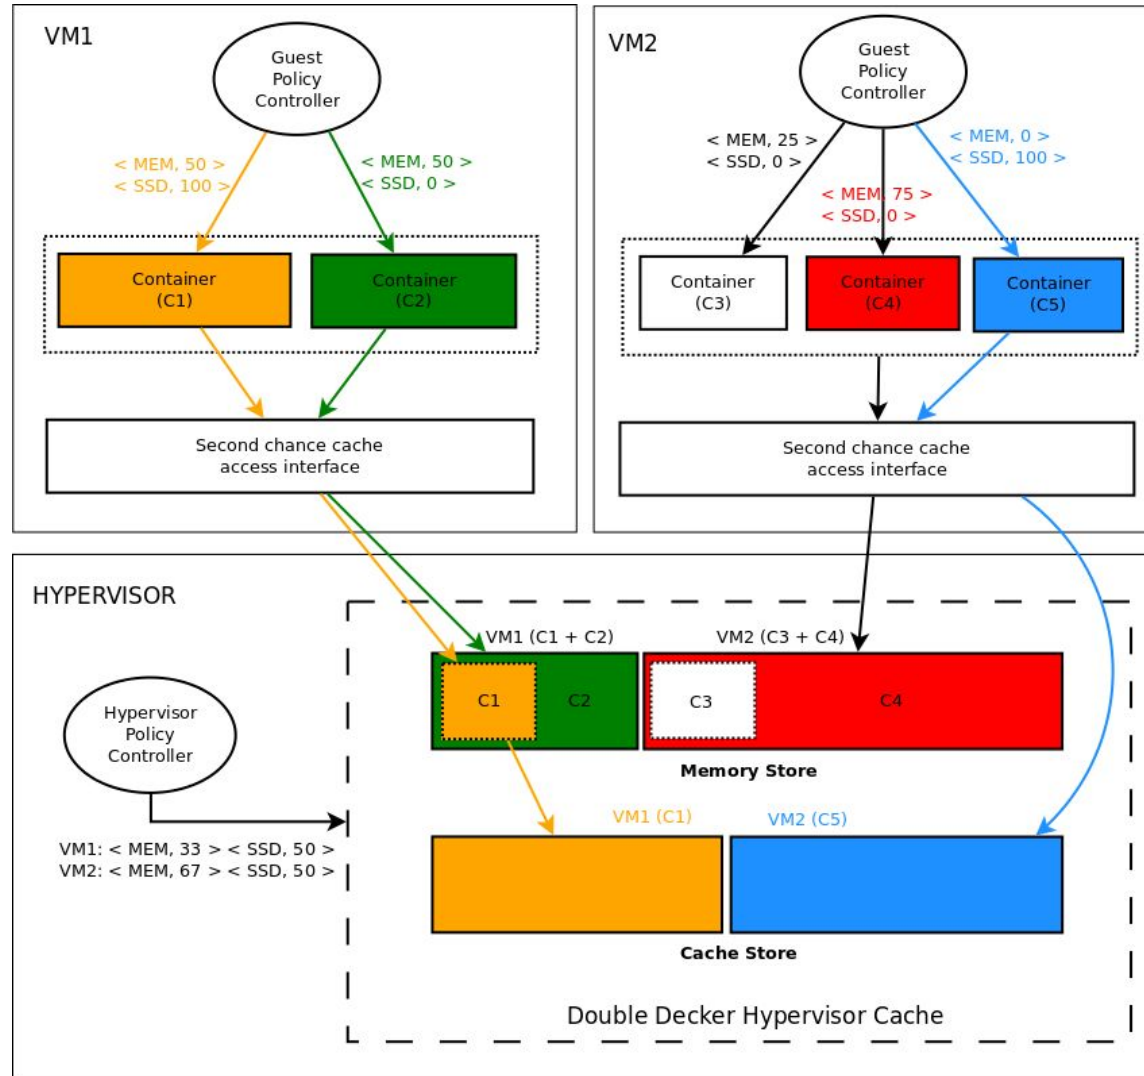


Fig-6: Implementation Details

Set of APIs to control caches

Set of APIs,

1. CREATE_CACHE
2. PUT_OBJECT
3. GET_OBJECT
4. DESTROY_CACHE
5. SET_WEIGHTS
6. EVICT_SPECIFIC_OBJECTS
7. EVICT_SET_OF_OBJECTS
8. MIGRATE_FROM_L1_TO_L2
9. MIGRATE_FROM_L2_TO_L1

Future Work

- ❑ Making use of developed APIs to map SLA policies into cache partitions
- ❑ Would hint passing from VM to Host help in cache partitioning ?

Thank You !

Any Questions ?